Claude Code Capabilities

Workshop 01

May 2026

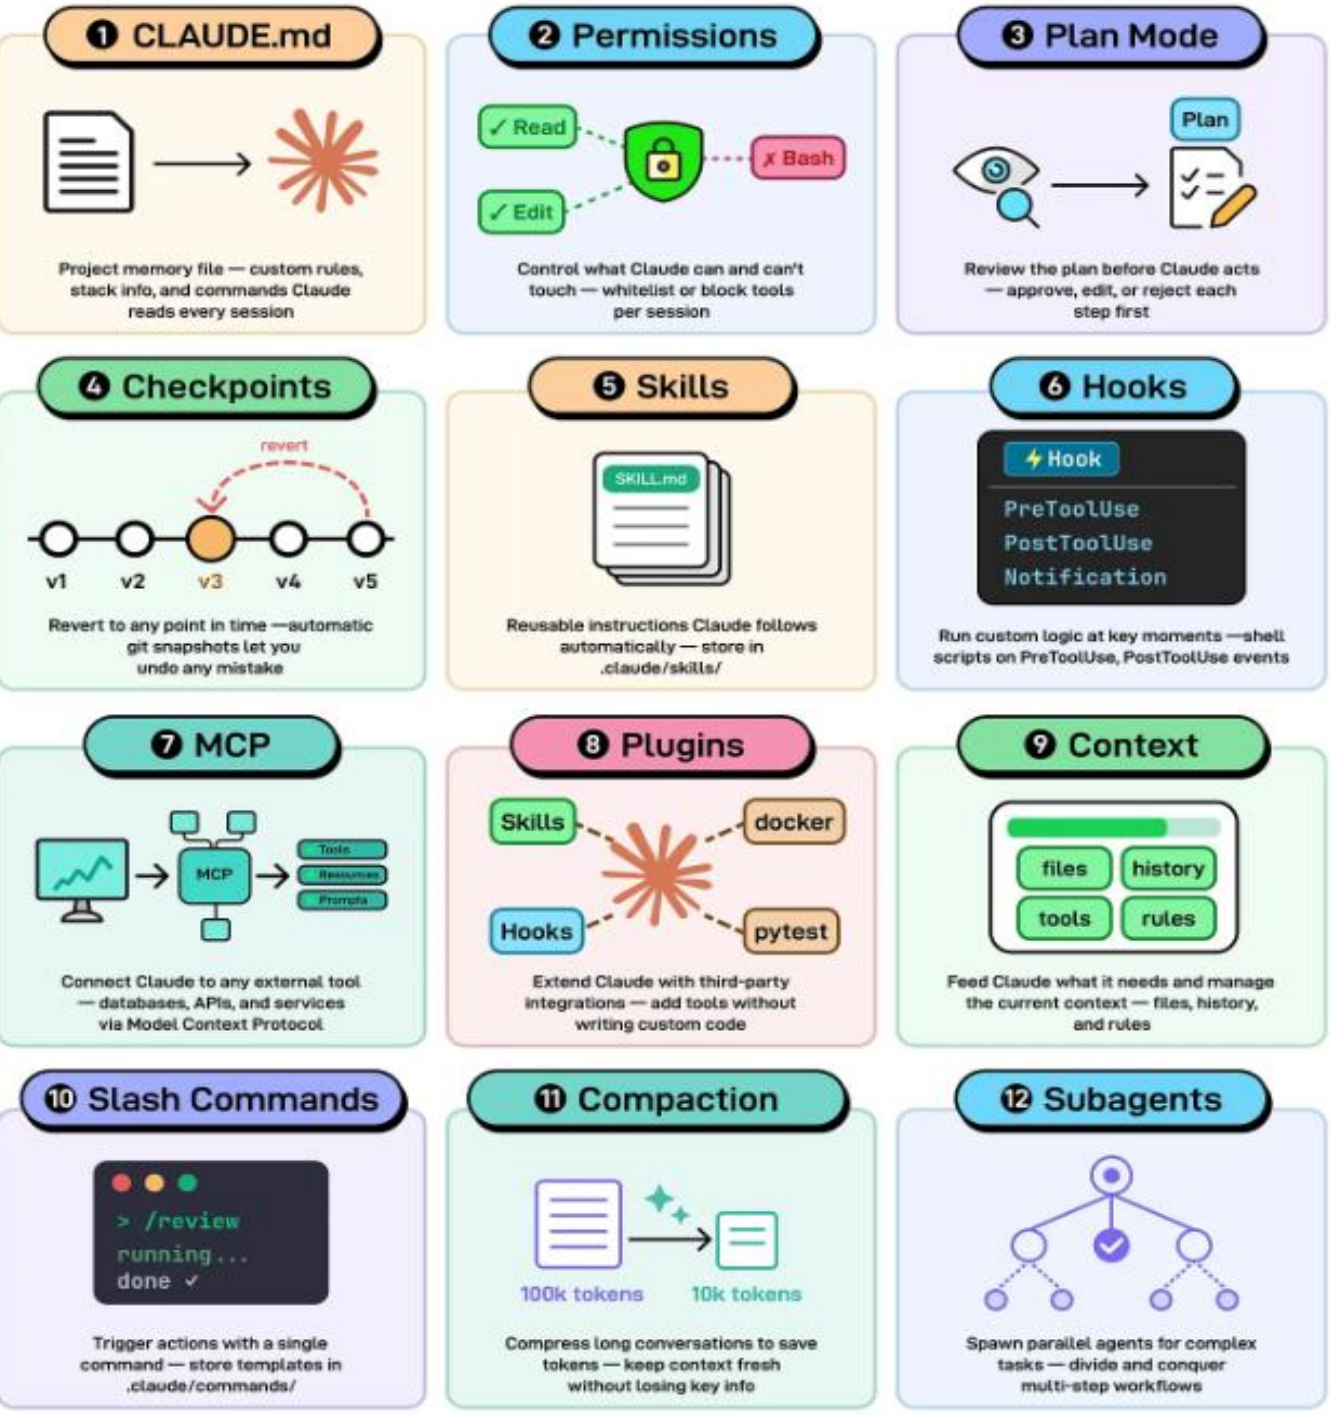

Topic-Wise Overview

- Review Topic Theory
- 3-5 min Hands-on Demo
- QnA

Workshop-wise topics

Workshop 01

1. Claude MD file
2. Claude Usage (Command Line, Plugin, Desktop)

Workshop 02

1. Claude Superpowers

Thank You

™

Appendix

Product Owners Feedback

All key users want costs insights

Most Valued Capabilities

Pointwise Evaluation
Multi-metric scoring (Correctness, Relevancy, etc.)

Ground Truth Generation
LLM-generated QA pairs from knowledge docs

Top Requested Features

Cost & Efficiency Insights	Token usage, cost analysis per run
Eval Insights & Explainability	Justifications, summaries, trends
Workflow Improvements	Scheduled runs, recurring evals

Key UX & Workflow Improvement Requests

- Search & filter test questions in results (not just scrolling)
- Agent-friendly UI names when selecting agents for eval
- Recurring scheduled evals with threshold-based alerts
- Auto-recommendations to improve agent based on eval results
- Ground truth from more formats (not just PDF) + SME flagging
- Tweak ground truth tone/POV (e.g. new employee perspective)
- Model comparison eval (e.g. Gemini 2.5 vs 3.0 on same agent)
- Support for non-Q&A eval types (recommendation engines, itineraries)

User Sentiment

"Overall I love it! Saves a lot of time and very easy to use."
— **Nina Dakoriya**

Stated Priority Order

1. Workflow Improvements
2. Insights & Explainability
3. Cost & Efficiency

Preliminary Sprint-by-Sprint Roadmap

GANTT · Q2 2026

13 weeks · Apr 1 – Jun 27 · Status shown per initiative

Initiative	S1 Apr 1–14	S2 Apr 15–28	S3 Apr 29–May 12	S4 May 13–26	S5 May 27–Jun 9	S6 Jun 10–27	Status
COPILOT / AGENT 365 RESEARCH · P1							
Capabilities Assessment & Strategy	2 wks						○ Planned
AGENT OBSERVABILITY & MONITORING · P2							
Prism Hub Non-Prod Pipeline	3 wks						• Active
OTel Instrument Google Agent + Azure Dashboard			8 wks				• Active
EVAL CONSOLE ENHANCEMENTS · P3							
UDX Agent Eval + Baseline Metrics	2 wks						○ Planned
Auto Summaries & Insights		3 wks					• Active
Token / Latency Metrics			2 wks				○ Planned
CI/CD & PRODUCTION READINESS · P4							○ Planned
Automate Eval Pipelines (non-prod + prod)				4 wks			○ Planned
Swagger / API Docs						1 wk	
AUTOMATED AGENT ONBOARDING · P5							
Automated Onboarding Flow						→	○ Planned

Prism O&E Squad: Key Deliverables Overview

Prism Gen AI Evaluation | Past 12 Months

1 Gen AI Evaluation Console

- Pointwise evaluation of agent response quality across HR, AdSales, Studio Finance, Legal, etc.
- LLM Judge metrics: Correctness, Relevancy, Conciseness, Coherence
- Enables product owners to pinpoint and tune agent accuracy gaps

Impact: Structured quality assessment of agent responses to ground truth user queries

2 AI Ground Truth Generation Service

- Auto-generates QA pairs from knowledge source docs via LLM
- Integrated into Evaluation Console as a standalone feature
- Excel export for offline review, filtering, and re-upload by product teams ensuring human-in-the-loop

Impact: Accelerates ground truth creation from a previously manual process

3 Agent Observability Hub

- End-to-end monitoring via OTel Collector routing traces to Azure
- Observability Hub will cater to: Google, Prism custom, and Snowflake agents
- MVP demoed with dummy HR agent traces; next steps: Instrument live agents

Impact: Real-time, scalable monitoring and alerting across all Prism and OOB agents, reducing mean time to resolution for agent incidents

4 Safety, Compliance & Justification Flows

- Safety evaluation flow: grading on Tone, Toxicity, PII Leakage, etc.
- Compliance scoring validates agent readiness for deployment
- Evaluation Justifier explains LLM grading rationale per pointwise metric

Impact: Ensures agents are safe, compliant, and evaluations are justifiable

2026 PRISM Roadmap

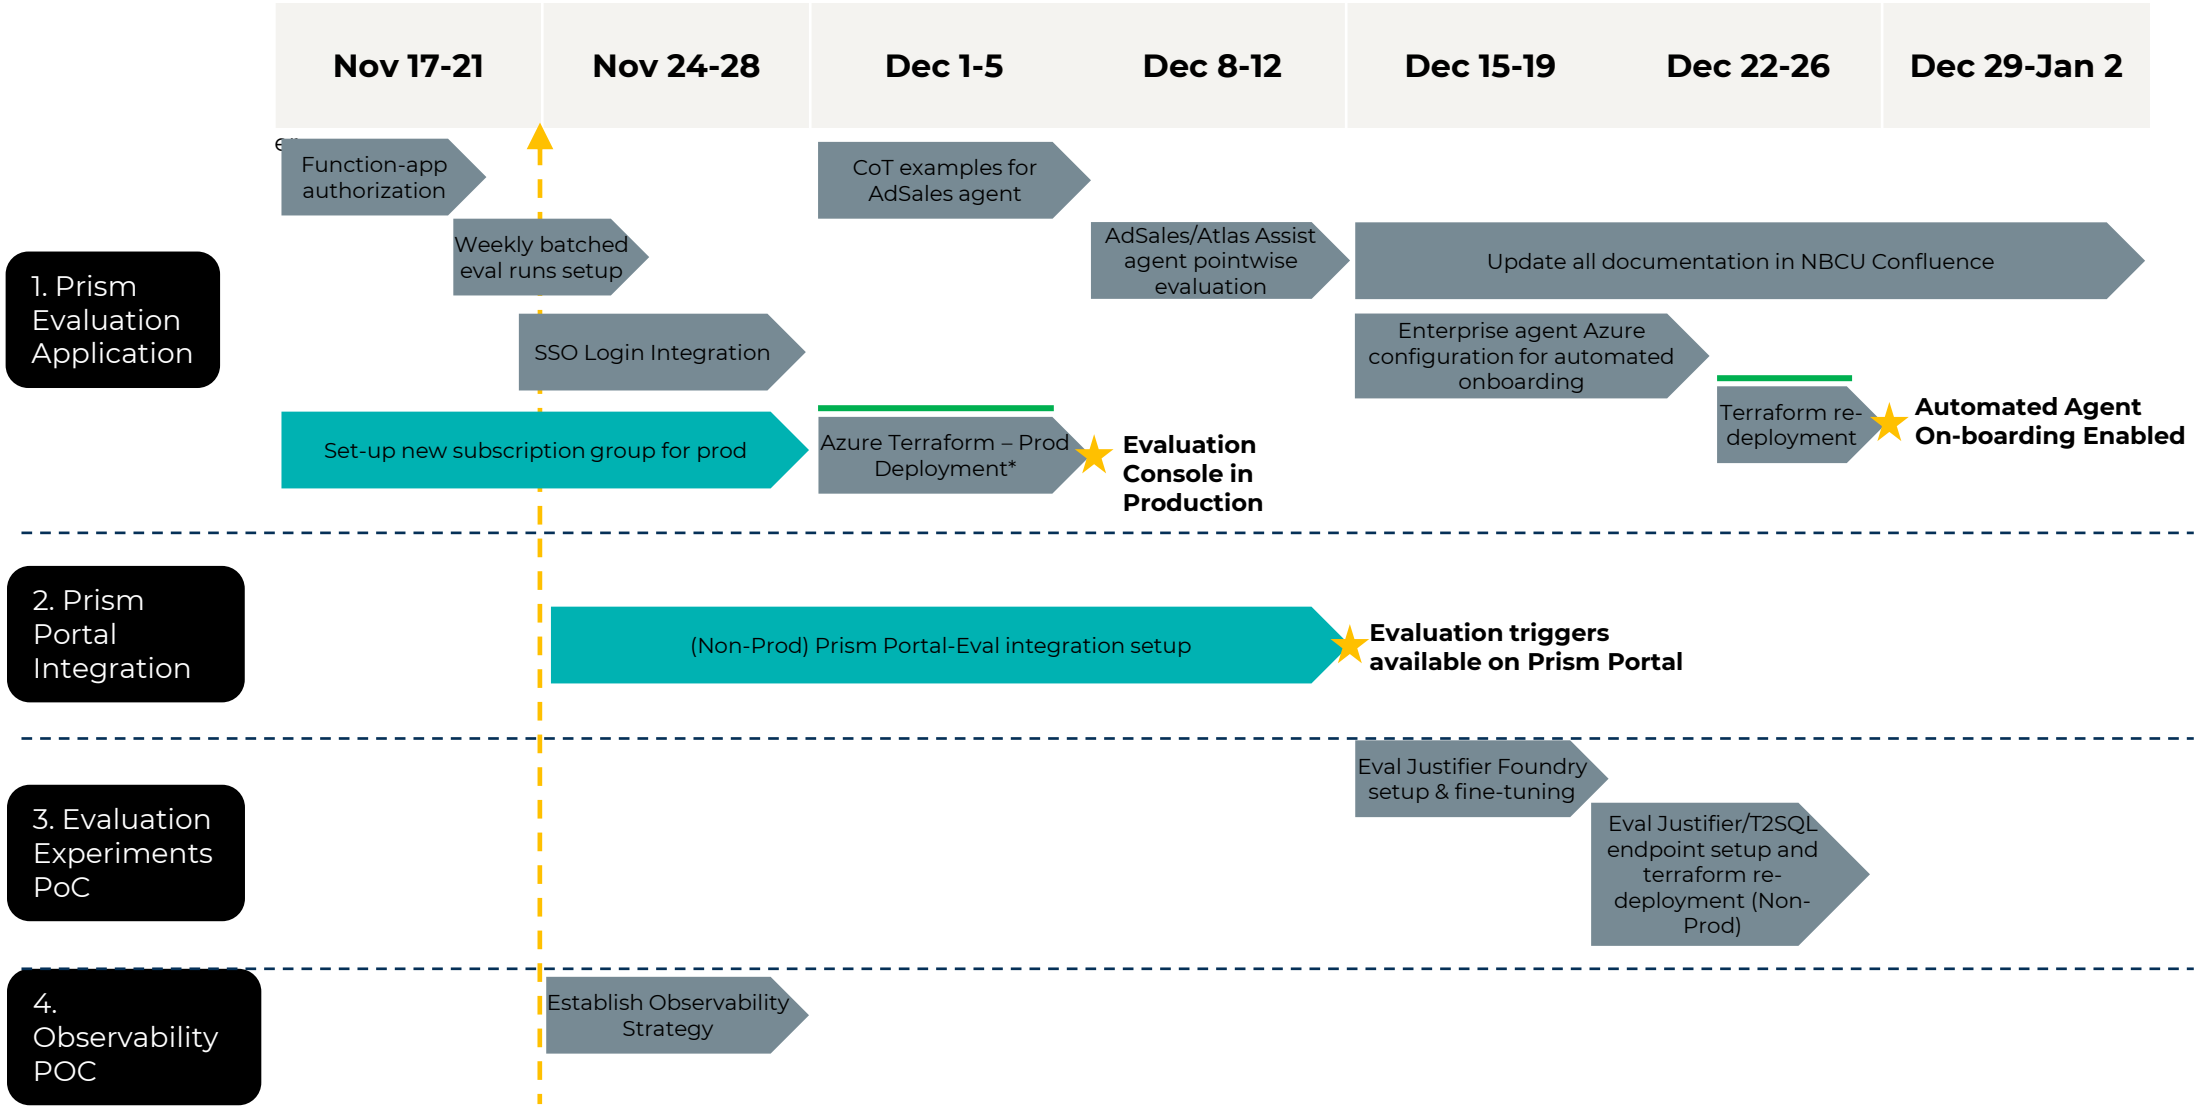

Prism Evaluation | E2E Eval Framework

A holistic evaluation framework which consist of multiple components integrated into one unified architecture

A **Expanding Evaluation Experiments**: A suite of agent evaluation experiments that can be selected in front-end UI for historical performance analysis and logging in DB

- Pointwise: As product testing quick performance of an agent.. I want to [text]
- Text-to-SQL: Evaluate the ability of agents to generate programmatically correct SQL code (AdSales Segment Hub)
- Pairwise: As product building two versions of an agent.. I want to [text]

B **Integration Into Prism Agent Self-Service**:

- Set-up integration of eval UI with Prism Portal for quick agent evaluations
- Automate seamless onboarding of new agents to existing eval experiments

C **Automated Github Actions Eval Regression Tests**: A series of evaluation tests that ensure any changes made to the underlying agent uphold its most recent baseline performance on a set of eval metrics.

D **Online Monitoring and Observability**: These features allow us to continuously trace and observe production agents for any output errors, prompt failures or hallucination spikes

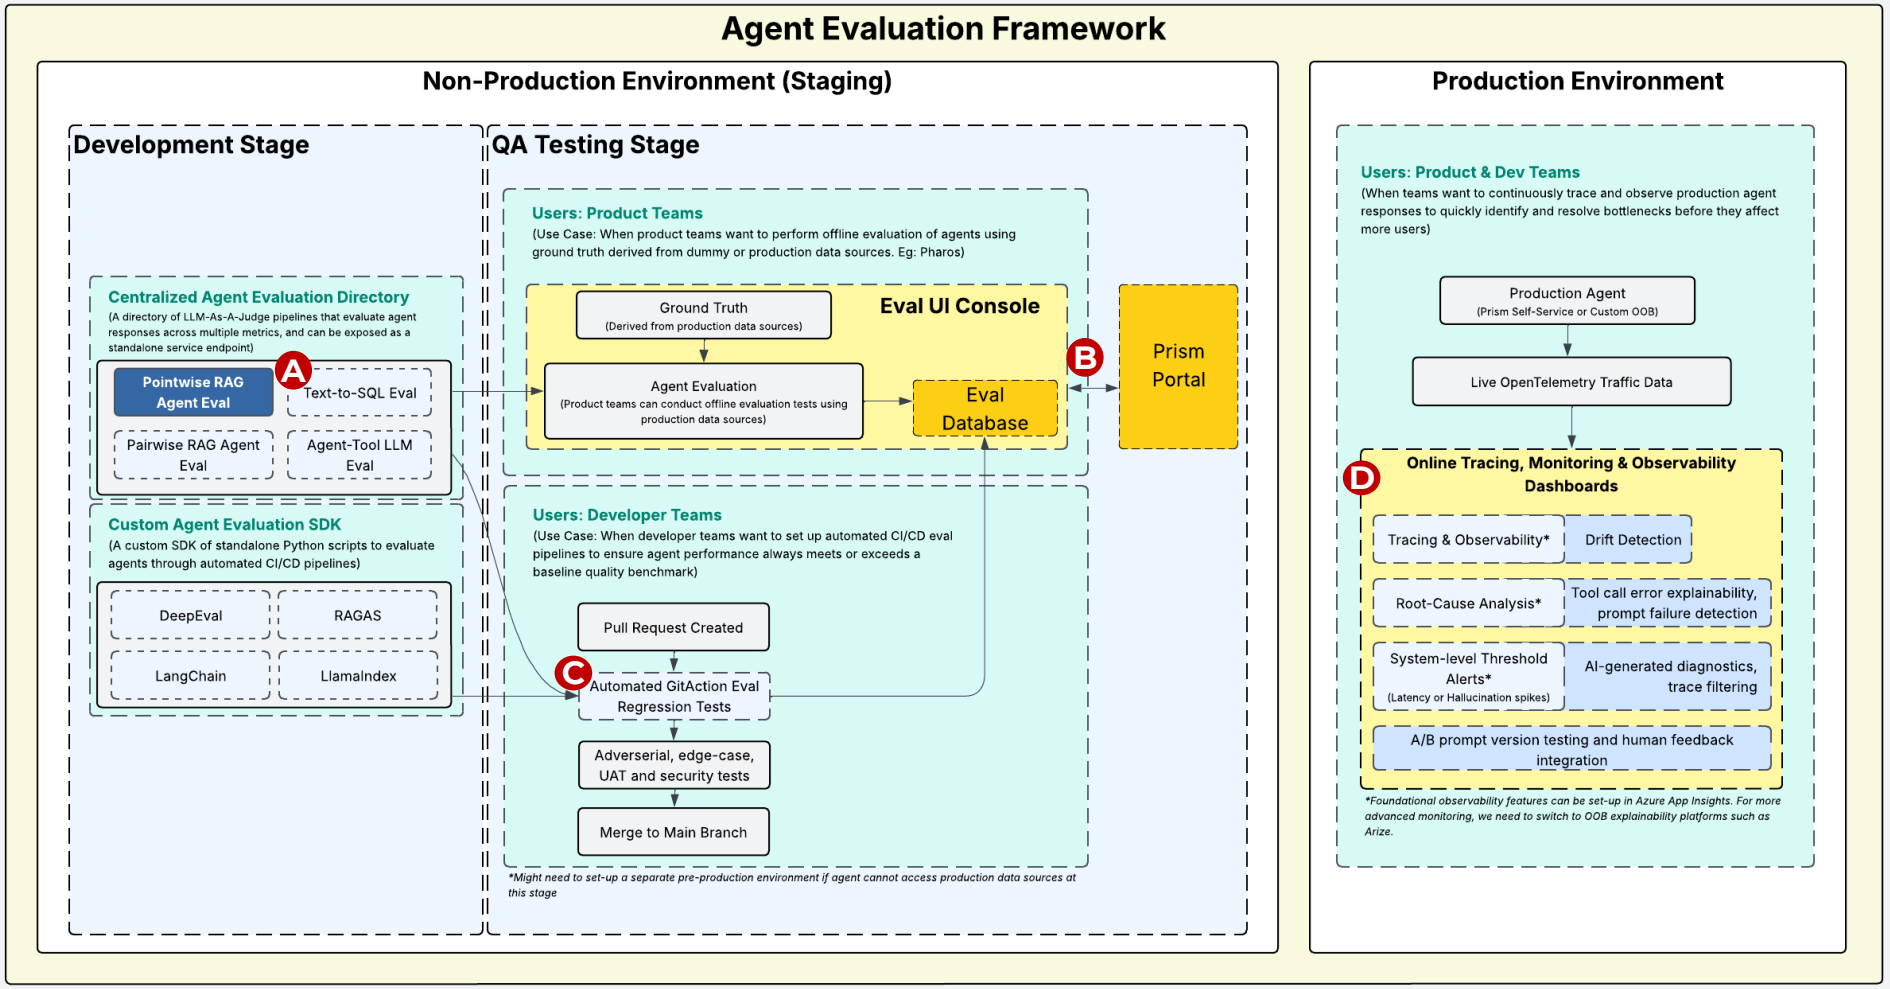

Supporting Business Impact Worksheet

PRISM O&E Business Impact Potential

IMPACT DIMENSION	Conservative	Moderate	Interactions/Evaluations Per Month		Cost Savings Per Evaluation/Interaction	Impact Range	
Evaluation Operational Efficiency	50%	80%	25000	100000	$100	$15,000,000	$96,000,000
			Value Agent Delivers Annually		Months of Incremental Value		
Faster Agent Delivery	30%	60%	$ 500,000	$ 2,000,000	3	$ 125,000	$ 500,000
			Averaged Fully Loaded Developer Cost		Number of Agentic Developers		
Lower Engineering Costs	20%	40%	$ 280,000		50	$ 2,800,000	$ 5,600,000
			Regulatory, Compliance, or Legal Incident Cost Range		Number of Avoided Incidents		
Higher Agent Trust & Reliability	25%	35%	$ 1,000,000	$ 10,000,000	2	$ 2,000,000	$ 7,000,000

Directional scenario modeling using conservative industry benchmarks applied to NBCU's enterprise scale. All figures are illustrative estimates for planning purposes. Validation through controlled pilot testing required before projecting actual business outcomes.

The average fully loaded annual cost of an Agentic AI Developer or AI Platform Engineer is $240,000 to $320,000 depending on geography and experience. (Source: Salary.com, Levels.fyi, Indeed)
*Avoiding one major AI Regulatory, Compliance, or Legal incident is often worth between $1M and $10M in cost avoidance

Meeting Objectives

3 Key Takeaways

1. Confirm the foundation we've built for AI progress and value creation
2. Align on Q2 Priorities
3. Share opportunities to accelerate and optimize value in 2026

3 Key Asks

1. Can you reaffirm the Long-Term Vision and ambition for PRISM?
2. How can we (Blend) help to continue momentum and champion adoption of what we're building?
3. What is objective feedback you can share related to our team's collaboration and performance?

What We Accomplished & Where We Are

Prism Gen AI Evaluation — Our Journey

Q1 2025 – Q4 2026+

From concept to production within a year —the evaluation console delivers self-service automated evaluation and agent onboarding to Product Owners, while we are establishing the baseline for near real-time production monitoring

Jan 2025

Kickoff

Began building pointwise evaluation with LLM Powered Judge metrics

Jul 2025

Evaluation Framework

Centralized evaluation endpoint for Developers for all Prism and non-Prism created AI agents

Jan 2026

Evaluation Console Deployed

Self-Service evaluation and GT generation platform for Product Owners

✓ Live

Feb 2026

Expanded Framework

All enterprise agents on-boarded, scheduled evaluations, automated Prism agent on-boarding

TODAY

Q1 2026

Observability & Safety Compliance

Agent production monitoring and Safety & Compliance Scoring

Active

Phase 1: Foundational

Phase 2: Expansion & Self-Service

Phase 3: Production Monitoring & Observability

Phase 4: Closed-Loop Optimization

Prism O&E: Key Q1 '26 Accomplishments & Insights

Prism Gen AI Evaluation | Q1

Key Accomplishments

- **Deployed Prism Eval Console to Production** with enabled AD group authorization, establishing a centralized and **self-service** evaluation platform for **Product Owners** for all Prism AI created AI agents
- On-boarded **5+ Enterprise AI Agents**: Ad-Sales Audience, Employee HR, UDX Pharos, Studio Finance, OLI 2.0 using **700+** curated ground truth questions
- **Centralized Ground Truth Datasets** across agents: Transformed scattered evaluation data from blob storage into a query-able database
- **Hardened Evaluation platform security and deployment readiness** with SSO integration, IP allowlisting, scheduler automation, and production environment configuration
- **Designed observability stack** using OTel, Azure Event Hubs, and App Insights — enabling end-to-end agent tracing across agent platform architectures

Lessons Learned/Insights

- Out-of-the-box **LLM judges reject more responses than human** reviewers would — calibrating strictness thresholds per agent domain was essential to avoid false negatives that erode product owner trust in the eval platform
- Few-shot examples only improve eval accuracy when they **span the full variety** of domain edge cases
- Having a **human-in-the-loop is necessary** to validate and tune LLM-generated ground truth question-agent response pairs from knowledge source documents for evaluation purposes
- **Lack of Retrieved context** is necessary to evaluate agent response hallucination
- **Cyber reviews:** We learned this the hard way: architecture changes to the eval pipeline triggered security review requirements
- Weekly syncs with NBCU teams helped us learn and adapt in time

CLAUDE.md — Project Memory for Claude Code

What Is CLAUDE.md?

A Markdown file loaded into every Claude Code session **before your first message**. It provides persistent project context across sessions.

Key Insight: Claude treats it as context, not enforced config. More specific & concise instructions = higher adherence.

Auto Memory: Claude also writes

3-Level Memory Hierarchy (Most Specific Wins)

1. **User-level** ~/.claude/CLAUDE.md — personal defaults across all projects (lowest priority)
2. **Project-level** ./CLAUDE.md or .claude/CLAUDE.md — team standards, committed to git
3. **Subdirectory-level** ./src/api/CLAUDE.md — scoped rules loaded on-demand (highest priority)

Use CLAUDE.local.md for personal prefs (gitignored). Use @imports to compose modular configs.

What to Include

- **Build/Test Commands** — exact scripts Claude should use (npm run test, make lint)
- **Architecture Map** — top-level directory layout, key modules, monorepo boundaries
- **Coding Standards** — only what linters can't enforce (naming, patterns, imports)
- **Workflows** — step-by-step procedures (e.g. creating API endpoints, PR reviews)
- **Domain Terminology** — map jargon/acronyms to code entities
- **Gotchas/Rules** — files never to modify, required validation, security constraints

Best Practices

- ✓ Keep under **200 lines** — longer files degrade instruction-following
- ✓ Use **/init** to scaffold, then curate — delete what's obvious
- ✓ Be specific: **"use 2-space indent"** not "format code nicely"
- ✓ Don't duplicate what linters handle — use **hooks** for auto-formatting
- ✓ Split large configs with **@imports** and **.claude/rules/** path-scoped files
- ✓ Iterate: add rules when Claude **repeats a mistake**, not preemptively

Demo Screenshots

Prism Gen AI Evaluation | Key Feature Demos

1 Gen AI Evaluation Console

Quality scoring of agent responses across business units

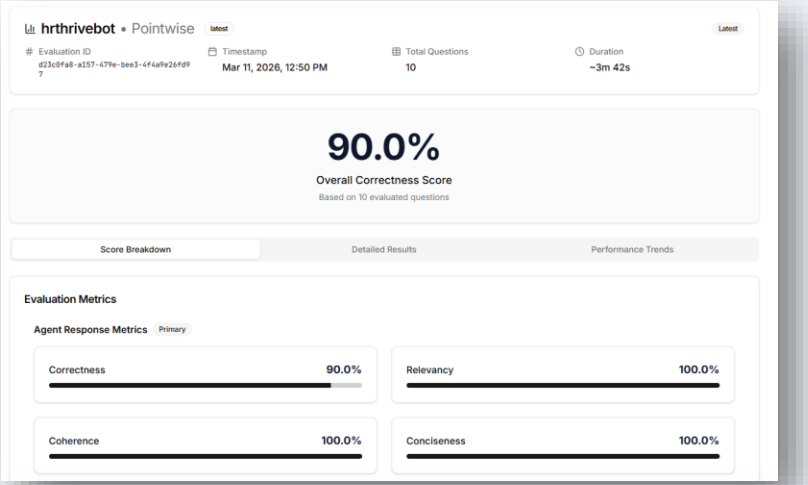

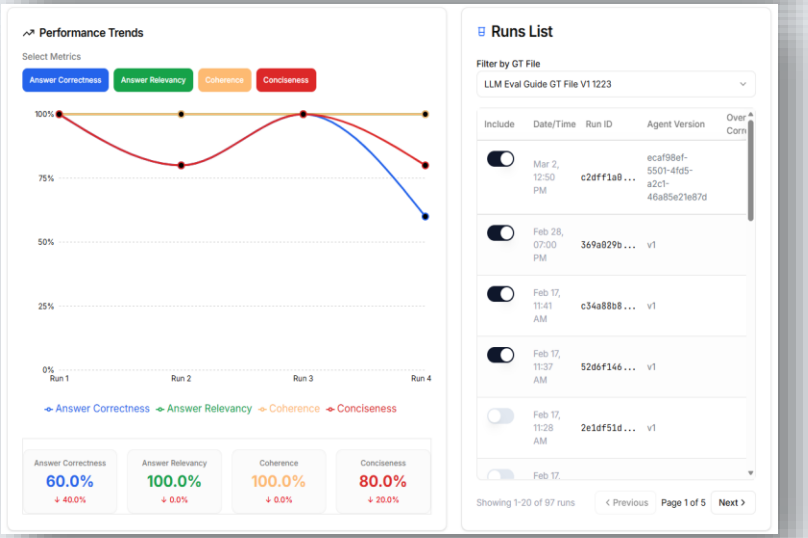

2 Agent Observability Hub

Real-time OTel monitoring across all Prism agents

3 Real-time Performance Alert System

Real-time alerts sent for anomaly Prism agent behavior

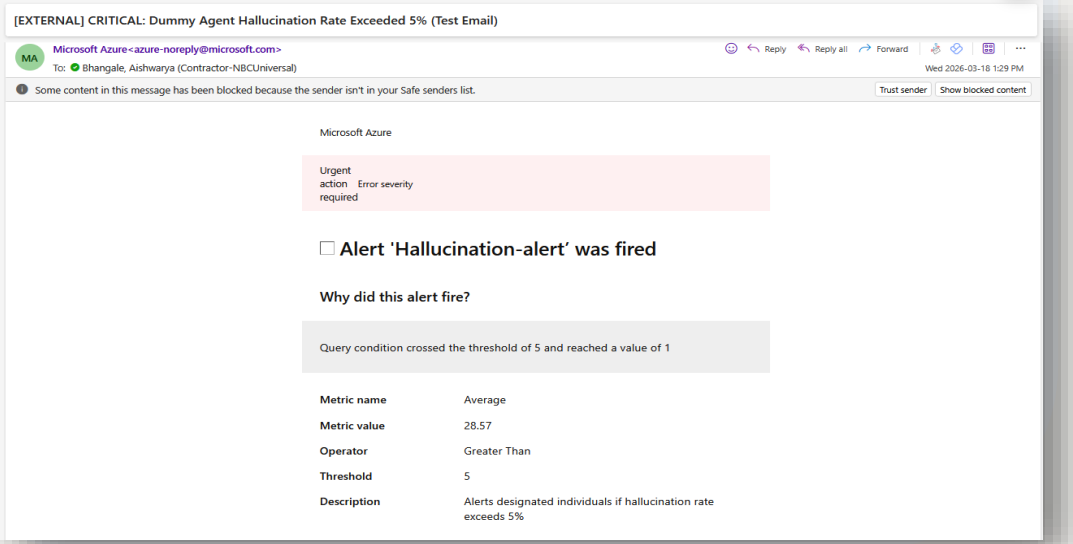

Business Impact

Business Value Enabled by O&E (Prism)

**In-flight*

Challenges Addressed

- Timely and inconsistent Agent Onboarding
- Manual Agentic Reviews
- Inconsistent Agent Performance Evaluation
- Custom logging and ad-hoc evaluation scripts
- Decentralized Agent Traceability

Core Capabilities Delivered

- Automated & Consistent Agent Onboarding*
- Accelerated Ground Truth Creation (from a previously manual process)
- Real-time, scalable monitoring and alerting across all Prism and OOB agents*
- Self-Service Offline Evaluation for Product Owners
- Online Evaluation using Open Telemetry*
- Consistent and scalable observability platform for any type of agent – Full view of Agent tracing signals and Agent performance*

Revenue Impact

- Faster Time to Value for Customer-facing & Business Analytics Agent Development and Deployment
- Higher Agent Reliability

Operational Efficiency

- Elimination/ significant reduction of manual Agentic Reviews
- Faster & Consistent Agent Onboarding
- Automated logging and evaluation scripts
- Lower Engineering cost per Agent
- Lower long-term TCO with one central Observability & Evaluation Platform

Strategic Enablement

- Platform and Tooling consolidation - Unified reporting & business intelligence
- Increased assurance of Regulatory Compliance and Data Governance
- Foundational Capability for Continuous Learning of Agent Behavior over time

Potential Business Value Impact of Observability & Evaluation

Directional Business Impact Enablement

Impact Dimension	Conservative	Moderate	Key Assumption(s)	Estimated Potential Annual Impact at Scale
Evaluation Operational Efficiency *enabling consistent automated evaluations (% cost reduction per review/evaluation)*	50%	80%	*Reduced Manual review of transcripts, conversations, Agent outputs, and decreased time to Ground Truth. $50 to $150 (Average of $100) cost per manual evaluation [Industry Benchmark].* *At Scale 25K to 100K interactions/evaluations per month*	**$15M - $96M** In Productivity Value due to reduced Evaluation Cost
Faster Agent Delivery *(% of efficiency increase)*	30%	60%	*30% to 60% decrease in Onboarding Time, Experimentation, Time to Production. Prior Avg. Time to Deploy = 6 months [Industry Benchmark]. Launching 3 months faster on average (range of 2 to 4 months) drives incremental business value per agent* *An Agent delivers $500K to $2M in annual business value [Industry Benchmark]*	**$125K-$500K** Per Deployed Agent
Lower Engineering Costs *(% efficiency increase)*	20%	40%	*20% to 40% decrease in Debugging, Rework due to missing Telemetry, Duplicate Tooling Effort, and Agentic Development* *Based on 50 Developers × $280K fully-loaded cost [industry benchmark] × 20-40% efficiency gain*	**$2.8M - $5.6M** In Productivity Value
Higher Agent Trust & Reliability *(% of leads correctly assigned)*	25%	35%	*Increased Explainability, Auditability resulting in less hard-to-explain decisions, inconsistent logs, and major regulatory or compliance violation** *Increased Task Rate, Reduced Hallucination Rate, Reduced Escalation Rate*	**$2M - $5M** Business Impact (Cost Avoidance or Revenue Uplift)
Platform & Tooling Consolidation *(% of TCO annual savings)*	$250K	$1M	*Reduces license, infrastructure & ongoing support costs – lowering tool sprawl & vendor overlap while driving consistent eval frameworks & metrics*	**$250K - $1M** In Annual TCO savings

Directional scenario modeling using conservative industry benchmarks applied to NBCU's enterprise scale. All figures are illustrative estimates for planning purposes. Validation through controlled pilot testing required before projecting actual business outcomes.

The average fully loaded annual cost of an Agentic AI Developer or AI Platform Engineer is $240,000 to $320,000 depending on geography and experience. (Source: Salary.com, Levels.fyi, Indeed)
*Avoiding one major AI Regulatory, Compliance, or Legal incident is often worth between $1M and $10M in cost avoidance

Q2 & Future Priorities

Prism AI Evaluation Platform: Roadmap Journey

Prism Gen AI Evaluation | Where We're Going

▸ Q2 2026 ROADMAP — Next Quarter

Evaluation Platform Rationalization
HIGH

Assess NBCU's preferred AI vendors (GAB, Snowflake, CoPilot Studio/Agent 365) to identify gaps and align evaluation platform strategy.

STRATEGIC IMPACT

Sets the architectural foundation for NBCU's long-term AI evaluation ecosystem — preventing costly technical debt and ensuring investments align with NBCU vendor strategy.

Observability Pilot
HIGH

Deploy Observability Hub in Non-Prod, instrument UDX UNI agent with OTel standard and Azure monitoring dashboards

STRATEGIC IMPACT

Foundation for unified Agent Observability Hub — establishing baseline ROI and potential for scaling to other NBCU AI agents.

UNI Agent on Eval Console
HIGH

Enable UDX agent evaluation on the Eval Console with verified baseline metrics for quality, correctness, and safety.

STRATEGIC IMPACT

Reduced time-to-production for the new UDX agent, ensuring consumer trust and establishing baseline quality and safety metrics before enterprise-wide rollout.

Production Readiness
MEDIUM

Automate evaluation pipelines across non-prod and prod for low-code platform agents, including quality gates.

STRATEGIC IMPACT

Embeds evaluation as a standard quality gate in NBCU's AI development lifecycle for low code agent— enabling safe, consistent AI deployments at scale.

Evaluation Console Enhancements
MEDIUM

Add automated evaluation summaries and insights, expanded safety & compliance workflows, and performance metrics (token/latency/cost)

STRATEGIC IMPACT

Enable business and product stakeholder to rapidly get the insights they need to make informed decisions about AI agent performance and trust

Automated Agent Onboarding
LOW

Streamline agent onboarding into eval — Q2 focus on Co-Pilot Studio agents.

STRATEGIC IMPACT

Eliminate manual registration overhead by auto-detecting and onboarding new Copilot Studios agents into the evaluation framework, reducing onboarding time from days to minutes.

⭐ Q3-Q4 VISION — Future Horizon

Unified Agent Observability

Unified Agent Observability Hub for centralized health, accuracy, and performance tracking across all Prism and non-Prism build AI agents.

STRATEGIC IMPACT

Enable continuous monitoring of production AI agents to detect performance degradation within hours, reducing incident response time and preventing user-facing failures

Automated Feedback Loop

Automated regression detection that triggers ground truth updates, suggested system prompt improvements based on real user

STRATEGIC IMPACT

Transform AI agent maintenance from reactive manual effort into proactive automated improvement—reducing Product Owner workload

Automated Agent Onboarding at Scale

Scalable onboarding across Google, Prism, Snowflake & Copilot w/ standardized evaluation and quality validation.

STRATEGIC IMPACT

Removes the last manual bottleneck in agent deployment — making Evaluation platform fully self-service at enterprise scale.

Prism AI Agent Governance Framework

Establish Responsible AI Agent gates thresholds ensuring consistent quality, safety & compliance at enterprise scale. Governance establishes policies, O&E enforces them.

STRATEGIC IMPACT

Transforms AI agent proliferation from a risk liability into a controlled growth asset—further democratizing agent creation without sacrificing safety, compliance, or accountability.

Adoption & Value Acceleration Paths

The PRISM O&E Value Flywheel

Agentic observability and governance are key to building trust in AI. Trust compounds over time. Each investment in PRSIM unlocks higher-value use cases, which generate measurable ROI, which funds further PRISM investment — creating a self-reinforcing cycle of enterprise differentiation.

Agentic AI Trust

High-confidence, from consistent and objective LLM evaluation, observability, and open telemetry governed across all platforms

Enterprise Edge

NBCU's PRISM platform becomes a structural competitive advantage

Reinvestment

ROI evidence funds roadmap expansion and governance maturity

AI Enablement Velocity

Acceleration of High-Impact Use Cases and Multi-Agent Orchestration

Visibility & Continuous Improvement

Unified observability across all agent platforms transforms operational data into a strategic feedback loop - continuously raising the bar on quality, reliability, and business impact

Measured ROI

Agentic Performance tracking, across key metrics ensure objectivity, and facilitates accurate ROI tracking

Considerations: Role-Based Evaluation & Aggregated Dashboard

Every agent interaction is logged, traced, and evaluated through our comprehensive observability framework. The Evaluation Agent runs continuously in the background — tracking performance, accuracy, and business impact.

Agent Performance

Accuracy, latency, throughput, cost per decision, and user satisfaction tracked in real time. Foundation metrics for AI performance.

Business Outcomes & Leadership Metrics

Impact measured across fill-rate improvements, and margin protection. Average dollar cost per agent run.

System Health

API uptime, data pipeline reliability, error rates, and model drift detection.

Continuous Learning

Feedback loops refine model routing, prompt quality, and agent behavior over time.

Risk & Compliance

Governance, policy violations, PII, and security relevant metrics.

EVALUATION DASHBOARD

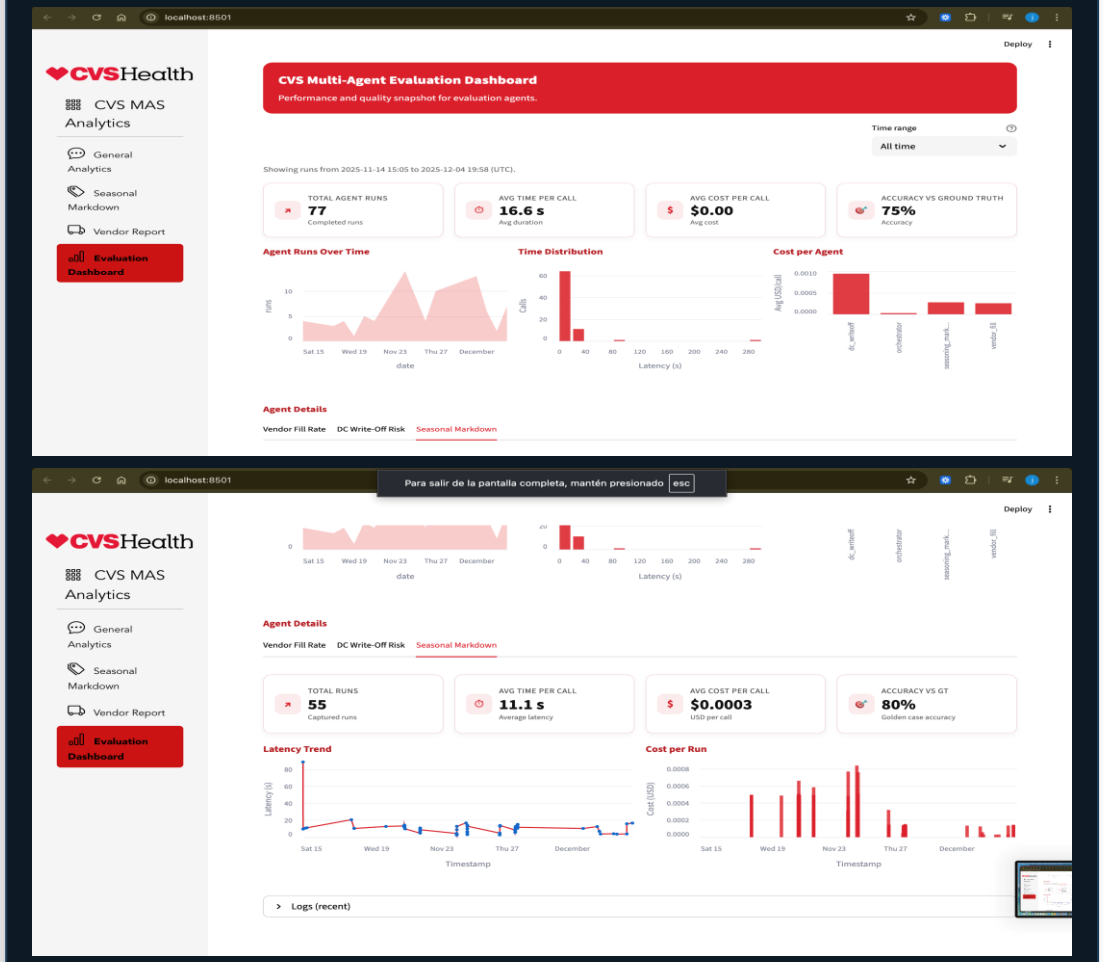

Role-Based Evaluation Dashboards

Different users see different views, all powered by the same telemetry and experiments. Each stakeholder gets the insights they need to make informed decisions about AI agent performance and trust.

Leadership

Trust scorecard with accuracy, SLA, risk, ROI trends, cost, time saved, and cost avoided metrics

Questions Answered: Can we trust the system, is it safe, is it delivering value, and are cost per transaction and throughput metrics scalable?

Product

Scenario pass/fail grids, output confidence score, disagreements, and annotation queues for QA

Questions Answered: Is the output correct, tone appropriate, complete, and aligned with expectations? Do we know when to include human review? where are the blind spots?

Engineers

Latency, error heatmaps, token drift, and trace replay for operational excellence

Questions Answered: Is the AI Agent's performance meeting throughput, accuracy, latency targets? Is the system healthy, stable, scalable and behaving as expected?

Risk & Compliance

Audit trails, PII handling, fairness and bias reports for governance. Assessment of Risk exposure

Questions Answered: Are we exposed to regulatory, legal, or reputational risk? Are we compliant with PII, and policies and can we audit it and prove it with trace coverage?

A New Way to Measure Success

Workflow_Success Token Score = (**L1_Score × α** + **L2_Score × β**) × **L3_Multiplier**

Layer 1	Layer 2	Layer 3
Execution Metrics (What AI Agent did)	**Validation Metrics** (What Humans found valuable)	**Business Impact** (What business outcomes achieved)

Measuring Success of AI Interactions

Success tokens to quantify AI impact by assigning measurable value to key outcomes, linking interactions directly to real business results.

Where:

- **α** = Weight of execution/Explicit signals (e.g., 0.4)
- **β** = Weight of validation/ Implicit signals (e.g., 0.6)
- **α + β = 1.0**

Driving Adoption to Capture Value

Adoption Discovery - Key Considerations to Work Through Together

Systems & Integration

- What downstream systems are highest priority for PRISM integration?
- Who owns the technical implementation of these integrations?'

Evaluation Use Case Prioritization

- Which performance evaluation business problems should PRISM's Evaluation & Observability team address next?
- What metrics define success for each use case?

Organizational Readiness

- Do teams understand what **PRISM's LLM Evaluation & Observability capabilities** are and why it matters?
- What training or enablement is needed before teams can fully lever Evaluation & Observability capabilities?
- What is the Communication plan to drive adoption?

Ongoing Evolution

- Consideration of aggregating performance data from PRISM Portal, Amplitude, and Agentic Platforms (Snowflake, Microsoft, GCP) into a holistic Executive Dashboard?

Meeting Objectives

3 Key Takeaways

1. Confirm the foundation we've built for AI progress and value creation
2. Align on Q2 Priorities
3. Share opportunities to accelerate and optimize value in 2026

3 Key Asks

1. Can you reaffirm the Long-Term Vision and ambition for PRISM?
2. How can we (Blend) help to continue momentum and champion adoption of what we're building?
3. What is objective feedback you can share related to our team's collaboration and performance?

Next Steps

1. Finalize Q2 Preliminary Roadmap
2. Finalize Q2 SOW

PRISM Gen AI Observability & Evaluation

Agentic AI Observability & Performance Tracking

Ensures agents are safe, compliant, and evaluations are justifiable

NBCUniversal